

Himavarsha Sreenivas

Houston, TX | himavarsha.2403@gmail.com | +1 346-397-0950 | [in LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Scientist with 3.5+ years of professional experience building predictive models, ML pipelines, and BI reporting solutions. Proficient in Python (scikit-learn, PyTorch, TensorFlow), SQL, and statistical analysis (hypothesis testing, A/B testing), with production experience in Power BI/DAX for executive dashboards. Hands-on project experience in NLP (transformer fine-tuning, text classification) and predictive analytics. Currently completing an M.S. in Engineering Data Science at the University of Houston.

EXPERIENCE

University of Houston Law Center - BI Developer / Data Analyst

Jan 2025 - Present

- As a Graduate Assistant, designed and built a Power BI financial reporting environment for a dean-level audience, including developing the underlying data model across 10+ source tables writing DAX measures and Power Query transformations to support drill-down reporting by department, account, and journal across 35 departments totaling \$34M+ in approved funding.
- Implemented row-level security (RLS) across multiple stakeholder roles, reducing manual data-scoping and access-review effort by ~30% and ensuring appropriate data scoping per department.
- Validated reported figures against official university financial statements line by line, improving reporting accuracy and reducing discrepancy-resolution time by ~25%

SHERP Scholar - Humana Institute, University of Houston

Jun 2026

- Selected for Humana SHERP (Summer Health Equity Research Program) 2026, an intensive interdisciplinary program at UH bringing together graduate students to address real-world public health challenges.
- As part of the Maternal Health team, led development and UI/UX design for Luna, a bilingual (EN/ES) companion app for high-risk pregnant and postpartum women in Harris County, featuring mood/symptom tracking, a personalized health score, local resource discovery, appointment prep tools, peer community, and IPV safety features.

Ernst & Young (EY-GDS) - Associate Analyst

Aug 2022 - Dec 2024

- Developed predictive and optimization models for lead scoring and customer segmentation using Python (NumPy, Keras, TensorFlow, Scikit-learn, SpaCy), improving forecast accuracy by 40%.
- Wrote custom SQL procedures in MySQL to clean, join, and aggregate multi-source datasets, reducing manual reporting effort by 30%.
- Applied hypothesis testing and A/B testing to validate client solutions, improving model deployment confidence and cutting post-deployment revisions by ~20%.
- Led a 4-member team onsite in Singapore to deliver a custom analytics dashboard under Agile sprints, managing last-minute client requests and delivered data-driven insights through visual presentations, accelerating stakeholder decision-making cycles by ~15%.

Data Science Intern - AiRobotica

Jun 2020 - Jul 2020

- Enhanced a client's ML algorithm and built forecasting models using Python (Scikit-learn, TensorFlow, Keras, Pandas, NumPy) on customer and retail market data processed in MySQL, improving model performance by ~15%.
- Conducted data quality assessments in Excel (VLOOKUP) and applied linear regression and ANOVA, identifying data relationships that improved analysis accuracy by ~10%.

PROJECTS

Electricity Demand Forecasting - AWS Cloud Deployment

- Built and deployed an end-to-end forecasting pipeline (ENTSO-E + OpenWeatherMap APIs, LightGBM model) as containerized services on AWS ECS/Fargate, with a Streamlit dashboard, EFS-backed persistent storage, and Secrets Manager/IAM for secure credential handling.
- Built a Financial Goal Tracker feature translating live forecasts into personal electricity cost estimates with configurable budget goals.

PulseBeat - SQL-Driven Music Recommendation & Analytics Platform

- Designed a normalized (3NF) MySQL schema with views, stored procedures, and triggers to power personalized recommendations, playlists, and ratings.
- Built a Streamlit + Python frontend with user login, recommendation engine, and an admin analytics dashboard.

AG News Topic Classification & Headline Generation

- Fine-tuned T5-base (223M params) for abstractive headline generation on AG News articles, benchmarking against DistilBART and PEGASUS across four decoding strategies.
- Achieved ROUGE-L 76.81% and BERTScore 94.68% on 120K+ news records.

EDUCATION

M.S., Engineering Data Science - University of Houston, Houston, TX - GPA: 3.89

Expected - Dec 2026

B.E., Computer Science and Engineering - Visvesvaraya Technological University, Bengaluru, India - GPA: 3.6

Aug 2018 - Jul 2022

HONORS AND CERTIFICATIONS

Microsoft PL-300: Power BI Data Analyst Associate – [View Credential](#)

DAAD RISE Professional Scholarship, Germany: among 58 selected worldwide - UH Recognition: [Facebook](#) / [Instagram](#) {declined}

EY Data Analytics & Data Visualization Bronze Badges ([Credly](#))

Salesforce Certified Administrator, Salesforce Certified Associate ([Trailblazer](#))

Additional: Architecting with Google Compute Engine Specialization, Google Cloud Ready Facilitation Program ([Google Cloud](#)), Python Programming (NSDC/ITM Edutech, Internshala)

SKILLS

BI & Reporting: Power BI (DAX, Power Query/M, RLS, semantic modeling), Tableau, Matplotlib, Seaborn, Plotly, dashboard design

Data & Databases: SQL, MySQL, Oracle, dimensional/star-schema modeling, Databricks, data architectures

Programming: Python (NumPy, pandas, scikit-learn, PyTorch, TensorFlow, Keras, spaCy, NLTK)

Machine Learning: Predictive modeling (ANN, SVM, Random Forest, XGBoost, AdaBoost, CatBoost), Clustering, PCA, Deep Learning, A/B testing

Cloud & Tools: GCP, AWS, Apache Spark, Jupyter Notebook, Google Colab, PyCharm, VS Code, Git/GitHub, JIRA, Agile/Scrum